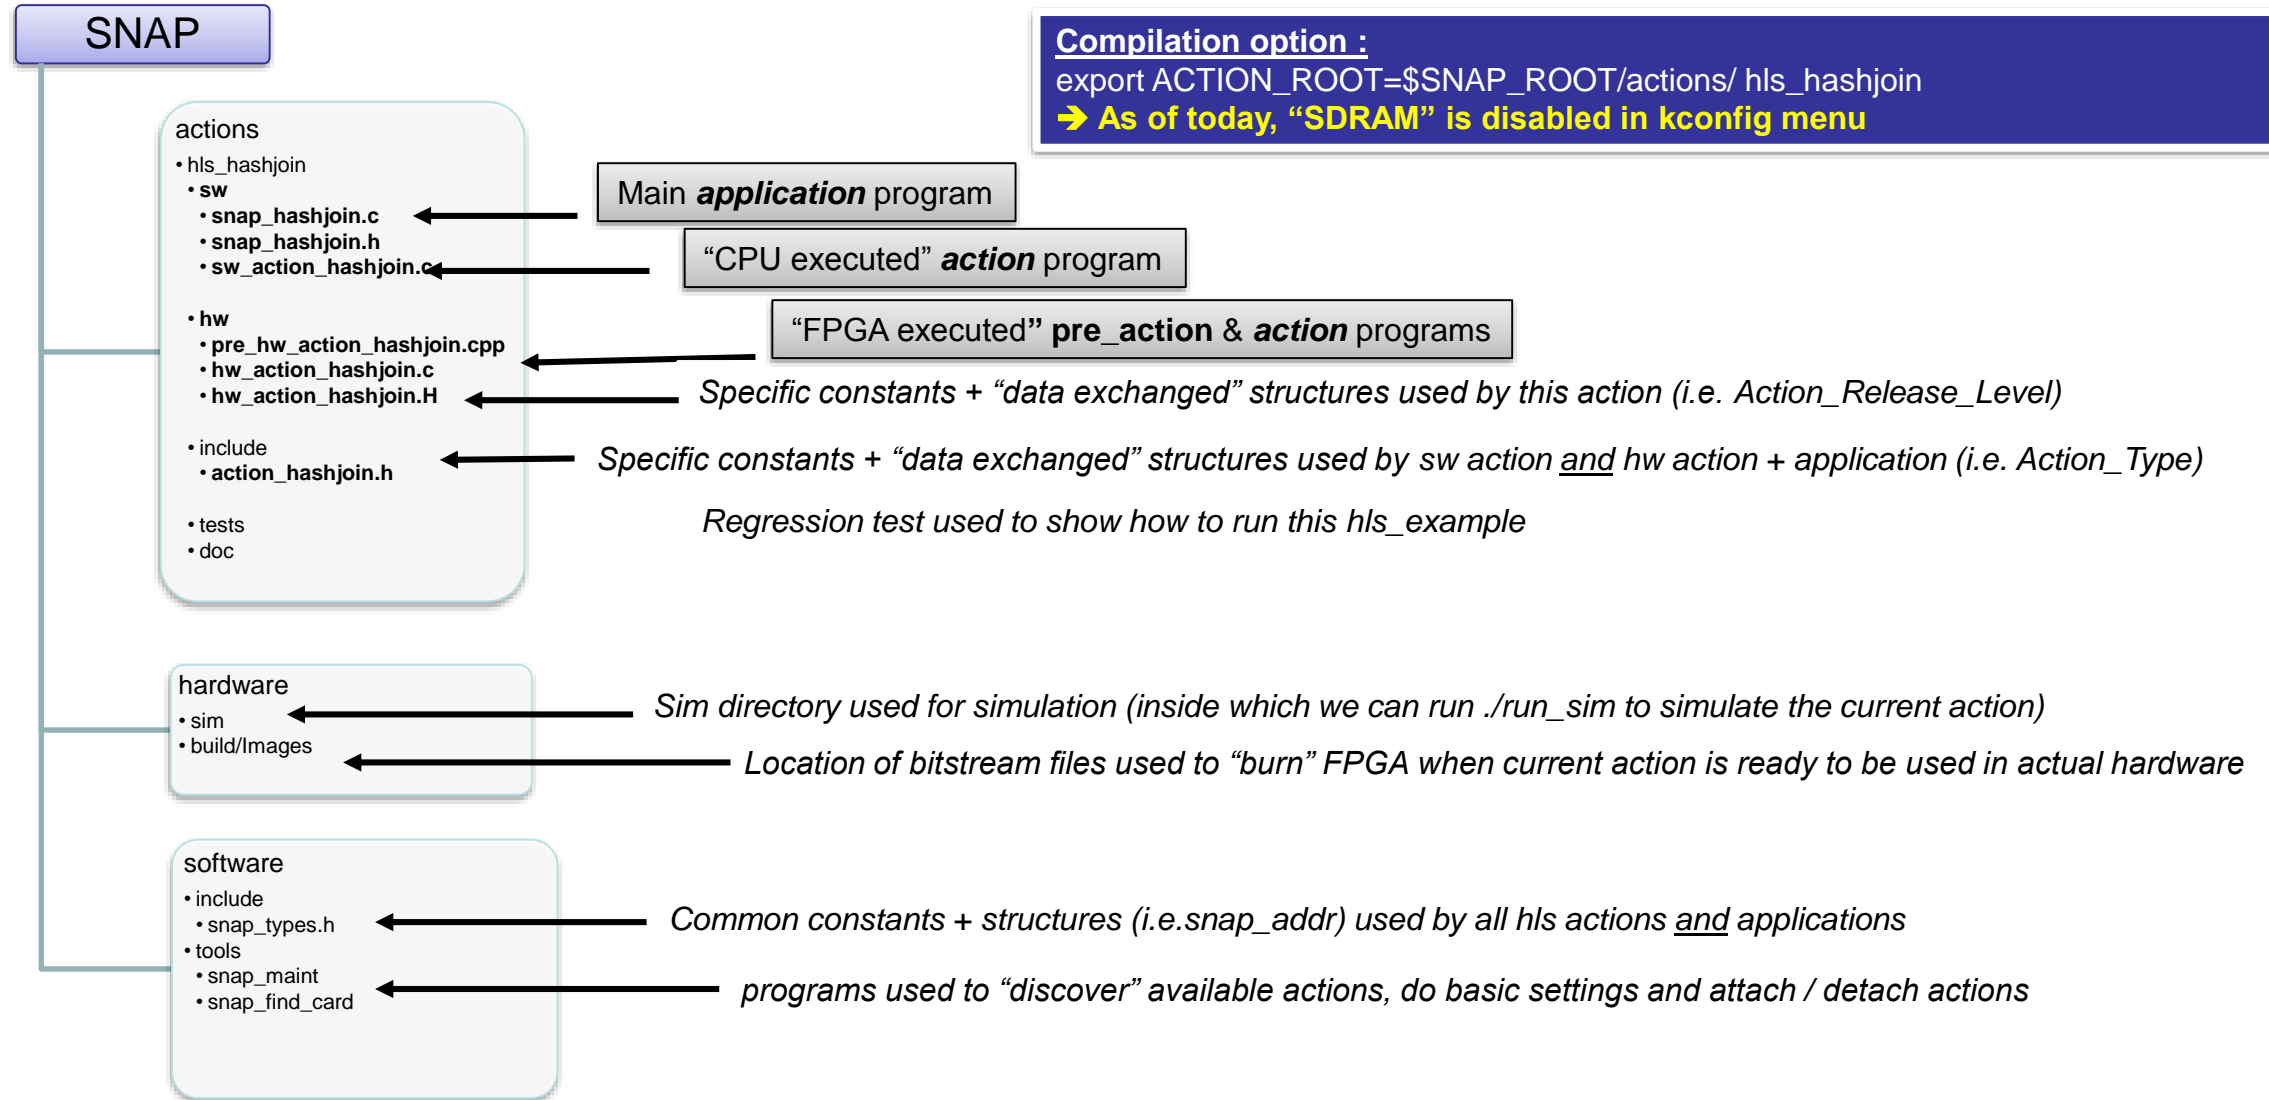


CAPI SNAP Education Series: User Guide

CAPI SNAP Education hls_hashjoin : howto? V2.2



Architecture of the SNAP git files



Action overview

Purpose: Port a hashjoin function

- Evaluate how tables can be managed with HLS
- Compare CPU and FPGA performances

When to use it:

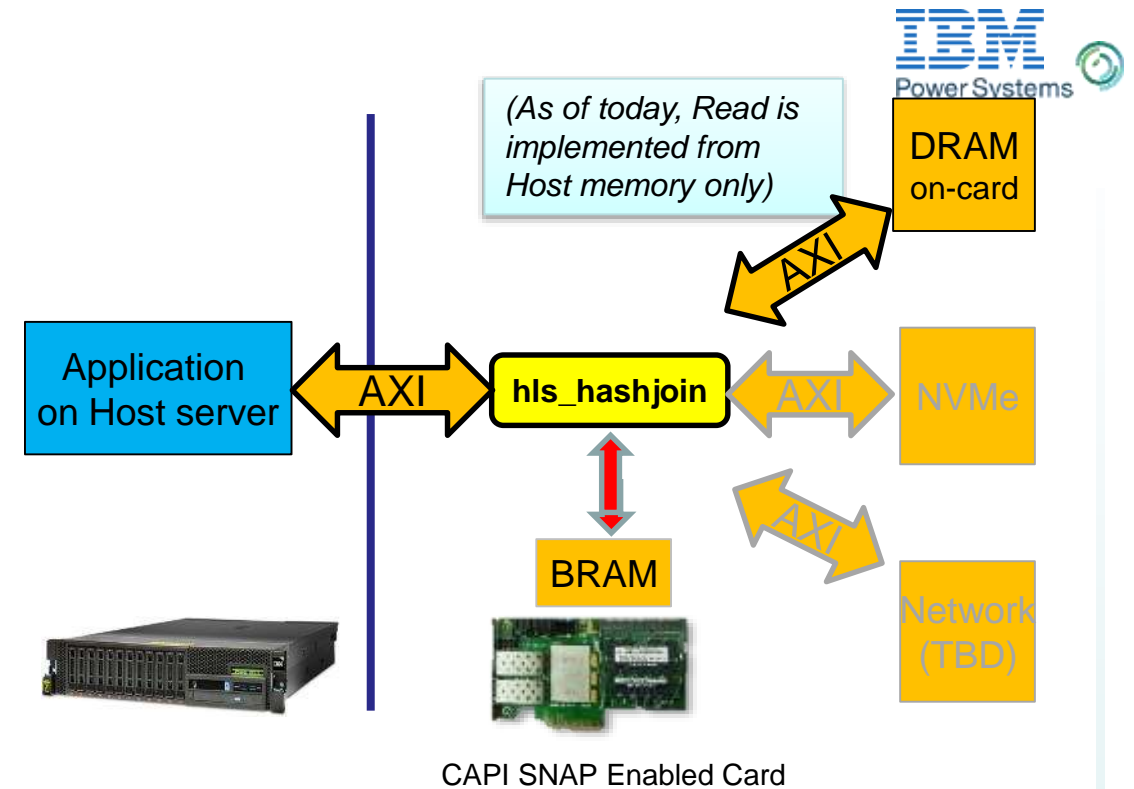
- Understand HLS constraints when working with large database

Memory management:

- All memory allocation is managed by the application

Known limitations:

- All data are 64 bytes aligned to ease access
- Data taken from Host memory instead of DDR



```
typedef struct table1_s {  
    hashkey_t name;      /* 64 bytes */  
    uint32_t age;        /* 4 bytes */  
    uint8_t reserved[60]; /* 60 bytes */  
} table1_t;
```

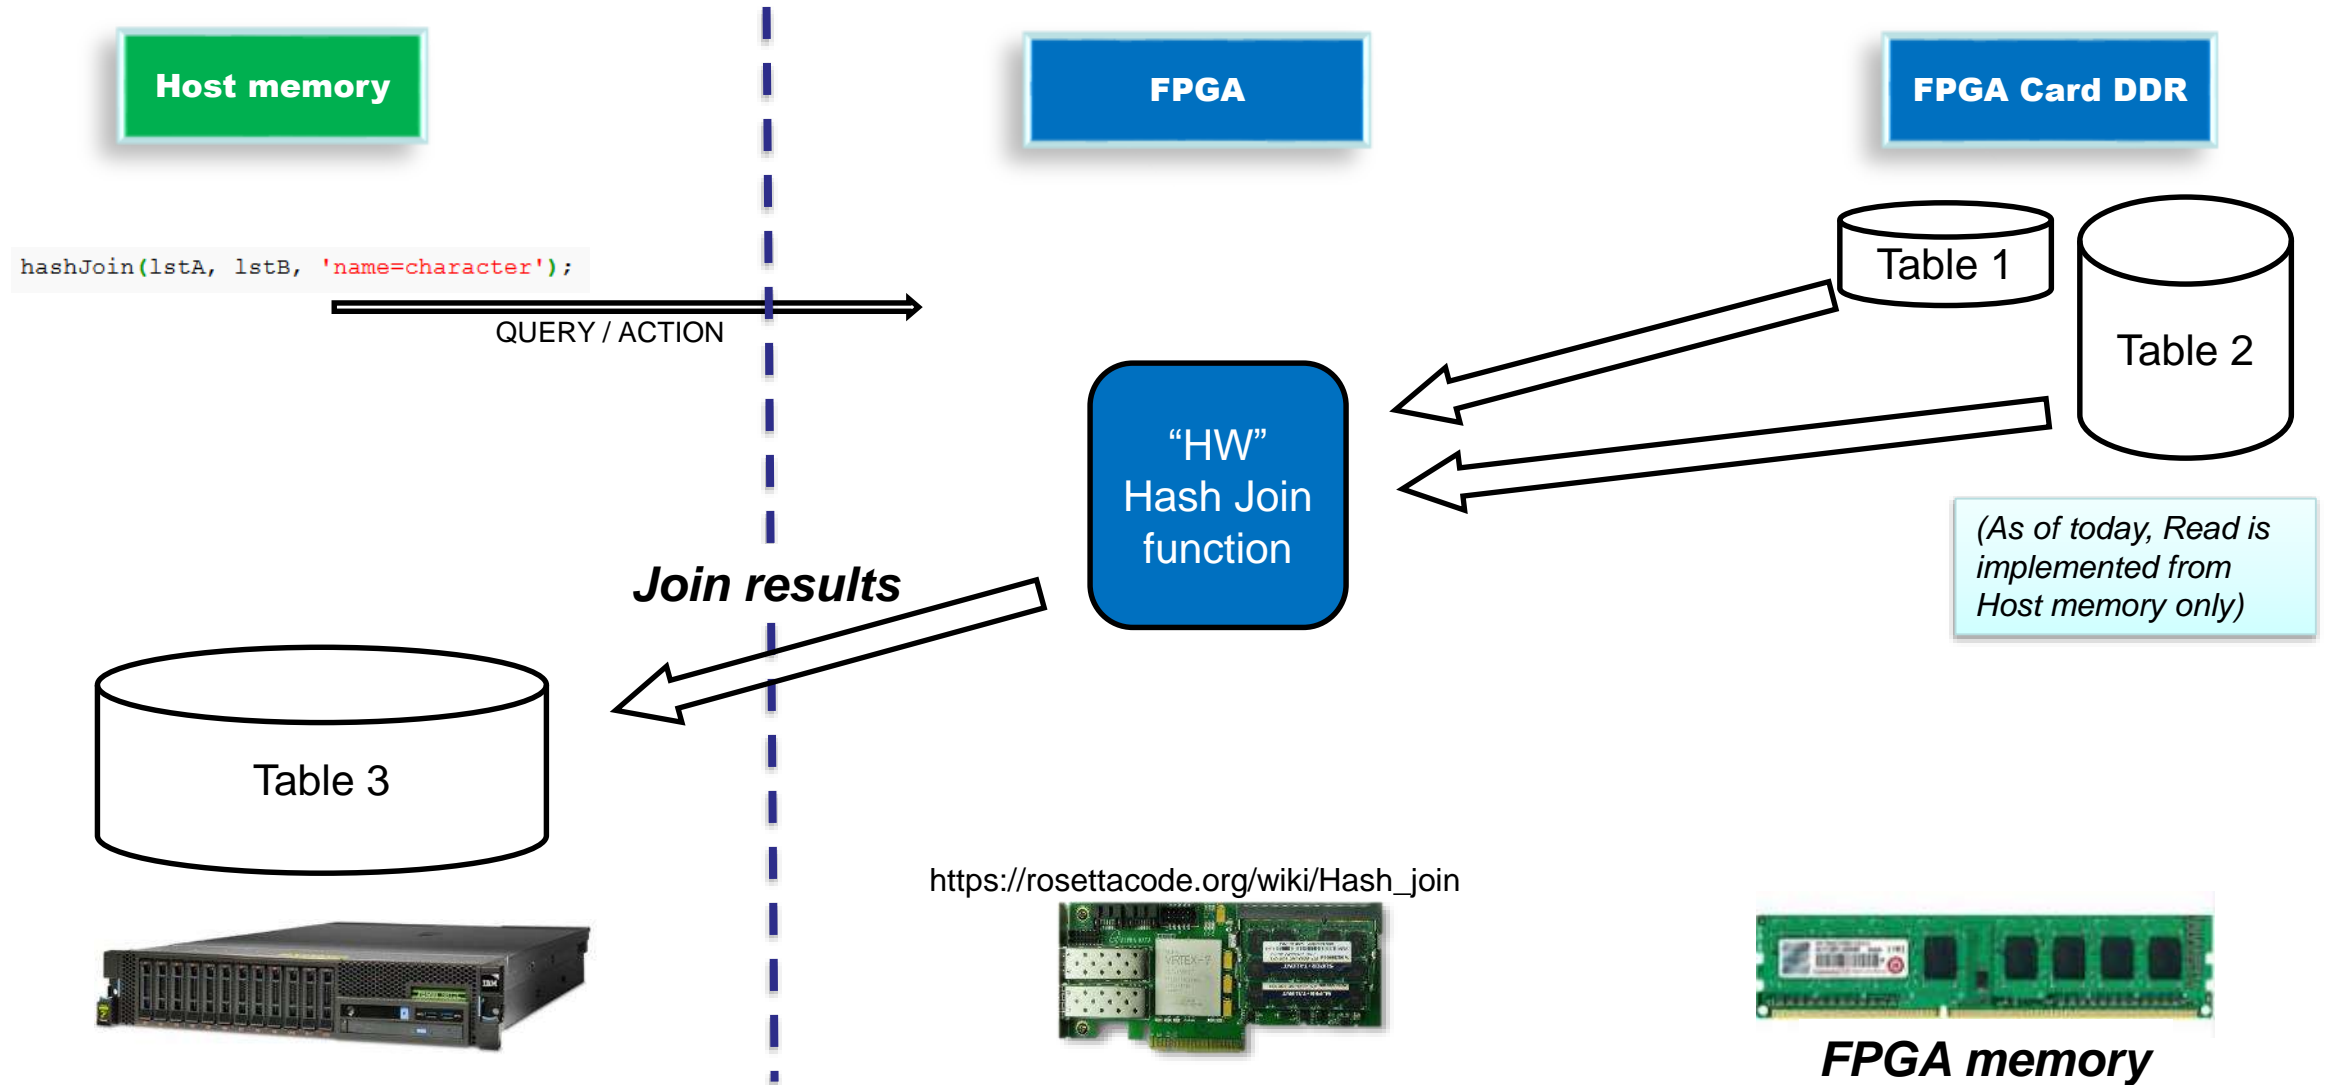
Hashjoin...an example

```
table1_t table1[] = {
  { .name = "Markus", .age=93 }      /* 0. */
  { .name = "Frank", .age=51 }      /* 1. */
  { .name = "George W.", .age=94 }  /* 2. */
  { .name = "Tercia", .age=63 }     /* 3. */
  { .name = "Secunda", .age=32 }    /* 4. */
  { .name = "Susanne", .age=99 }    /* 5. */
  { .name = "Tercia", .age=37 }     /* 6. */
  { .name = "Thomas", .age=71 }     /* 7. */
  { .name = "Joerg-Stephan", .age=89 } /* 8. */
  { .name = "Lisa", .age=47 }       /* 9. */
  { .name = "Julius", .age=75 }     /* 10. */
  { .name = "Glory", .age=83 }      /* 11. */
  { .name = "Melanie", .age=24 }    /* 12. */
  { .name = "Quintus", .age=77 }    /* 13. */
  { .name = "Prima", .age=52 }      /* 14. */
  { .name = "Andreas", .age=12 }    /* 15. */
  { .name = "Tercitus", .age=39 }   /* 16. */
  { .name = "Anders", .age=51 }     /* 17. */
  { .name = "Alexander", .age=38 }  /* 18. */
  { .name = "Dieter", .age=57 }     /* 19. */
  { .name = "Susanne", .age=48 }    /* 20. */
  { .name = "Melanie", .age=44 }    /* 21. */
  { .name = "Uwe", .age=50 }        /* 22. */
  { .name = "Jonah", .age=16 }      /* 23. */
  { .name = "Septus", .age=20 }     /* 24. */
}; /* table1_idx=25
```

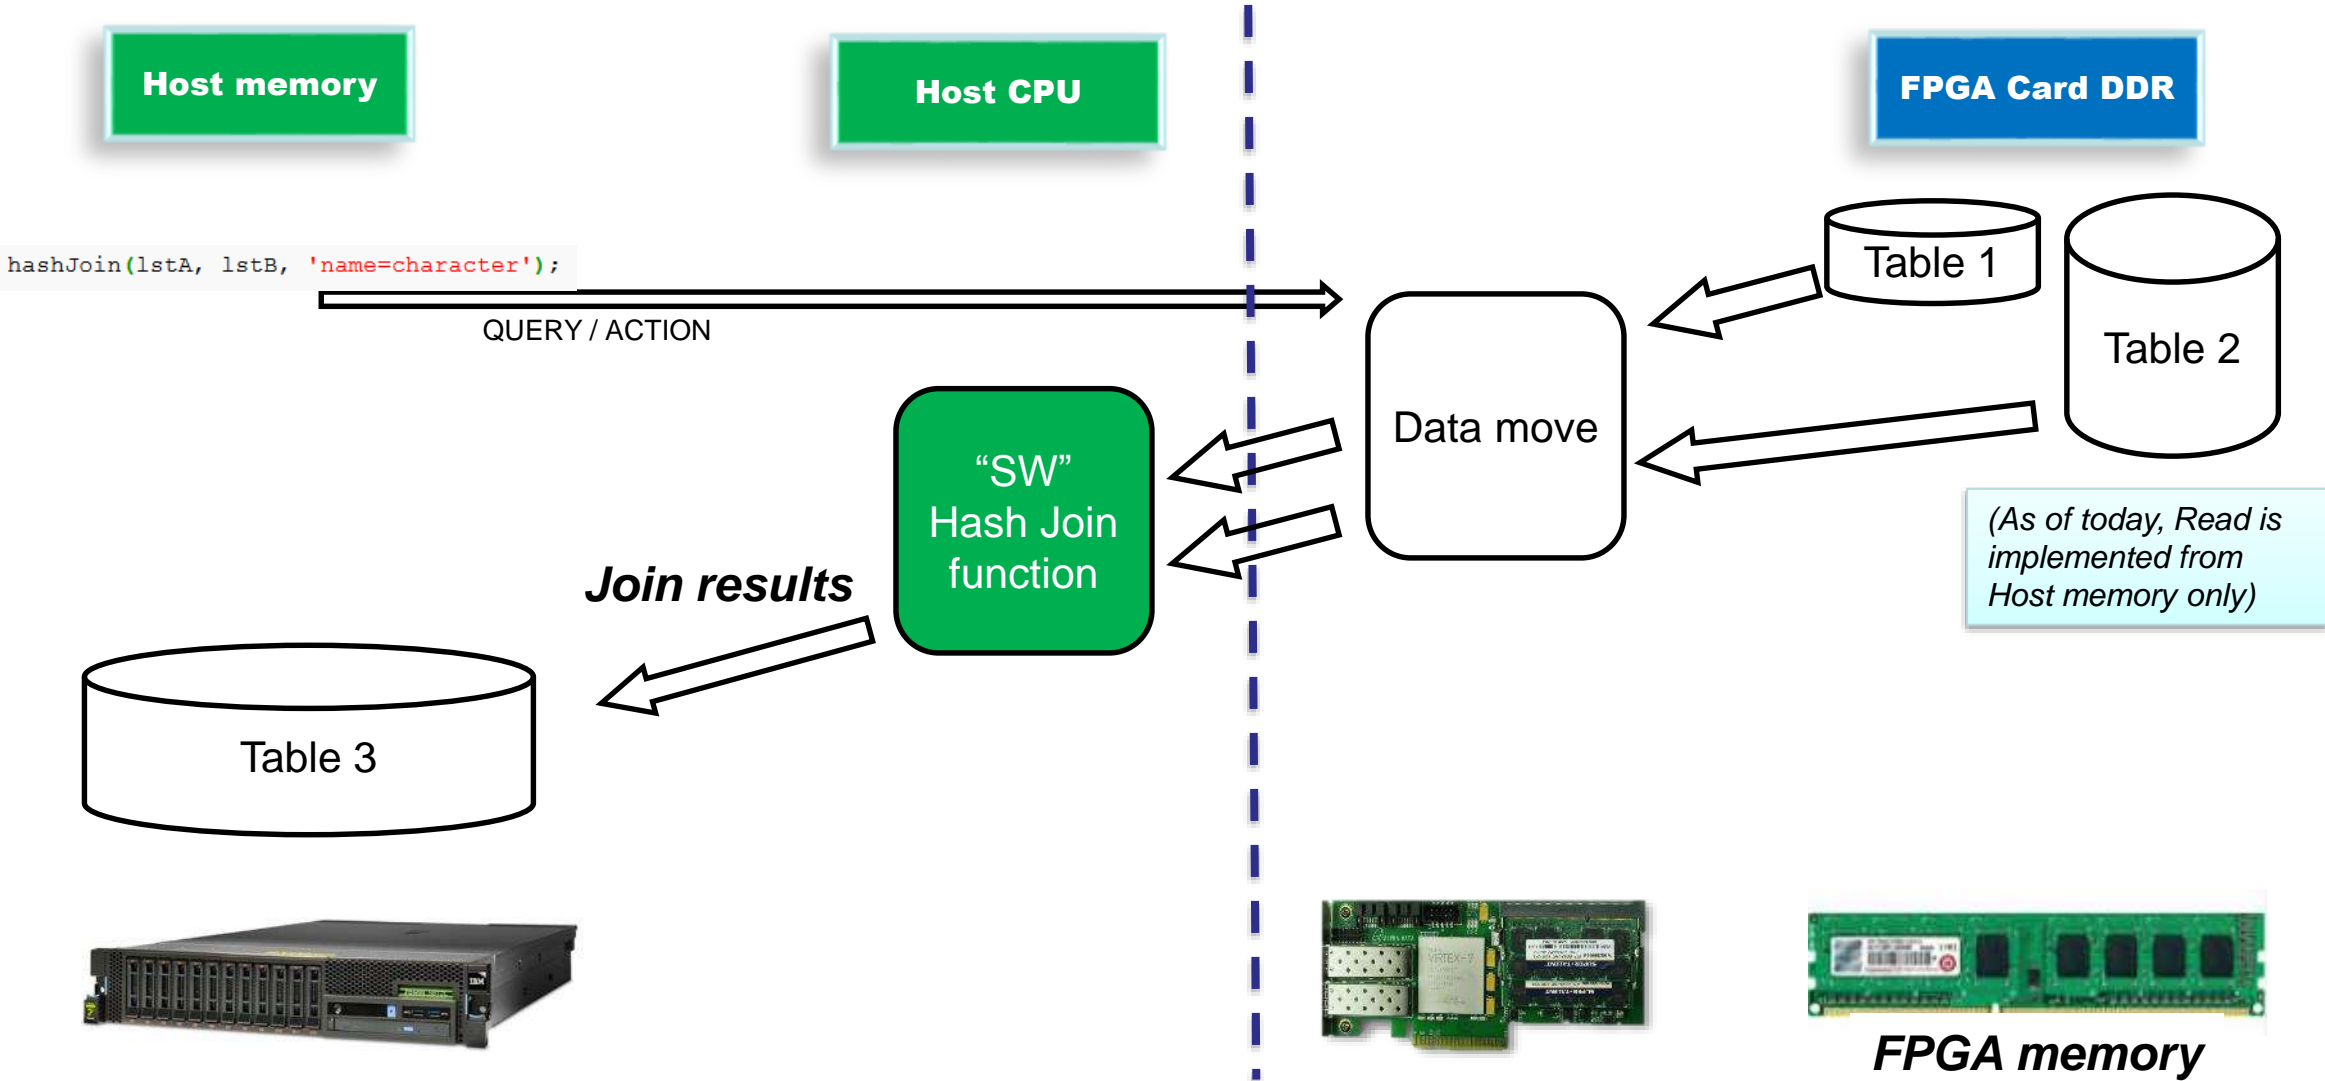
```
table2_t table2[] = {
  { .name = "Dirk", .animal = "Gorilla" } /* 0. */
  { .name = "Jonah", .animal = "Cat" }    /* 1. */
  { .name = "Horst", .animal = "Eagle" }  /* 2. */
  { .name = "Eberhard", .animal = "Dog" } /* 3. */
  { .name = "Eberhard", .animal = "Elephant" } /* 4. */
  { .name = "Quintus", .animal = "Greyling" } /* 5. */
  { .name = "Septa", .animal = "Gorilla" } /* 6. */
  { .name = "Mike", .animal = "Pike" }    /* 7. */
  { .name = "Maik", .animal = "Eagle" }   /* 8. */
  { .name = "George W.", .animal = "Cat" } /* 9. */
  { .name = "Septus", .animal = "Goose" } /* 10. */
  { .name = "Andrea", .animal = "Ghost" } /* 11. */
  { .name = "Susanne", .animal = "Antilope" } /* 12. */
  { .name = "Glory", .animal = "Trout" }  /* 13. */
  { .name = "Septa", .animal = "Dog" }    /* 14. */
  { .name = "Prima", .animal = "Cat" }    /* 15. */
  { .name = "Quintus", .animal = "Antilope" } /* 16. */
  { .name = "Mike", .animal = "Elephant" } /* 17. */
  { .name = "Primus", .animal = "Goose" } /* 18. */
  { .name = "Lisa", .animal = "Panther" } /* 19. */
  { .name = "Glory", .animal = "Gepard" } /* 20. */
  { .name = "Bruno", .animal = "Dog" }    /* 21. */
  { .name = "Septa", .animal = "Antilope" } /* 22. */
}; /* table2_idx=23
```

```
table3_t table3[] = {
  { .name = "Jonah", .animal = "Cat", .age=16 } /* 0. */
  { .name = "Quintus", .animal = "Greyling", .age=77 } /* 1. */
  { .name = "George W.", .animal = "Cat", .age=94 } /* 2. */
  { .name = "Septus", .animal = "Goose", .age=20 } /* 3. */
  { .name = "Susanne", .animal = "Antilope", .age=99 } /* 4. */
  { .name = "Susanne", .animal = "Antilope", .age=48 } /* 5. */
  { .name = "Glory", .animal = "Trout", .age=83 } /* 6. */
  { .name = "Prima", .animal = "Cat", .age=52 } /* 7. */
  { .name = "Quintus", .animal = "Antilope", .age=77 } /* 8. */
  { .name = "Lisa", .animal = "Panther", .age=47 } /* 9. */
  { .name = "Glory", .animal = "Gepard", .age=83 } /* 10. */
}; /* table3_idx=11
```

Hash Join : data are in card DDR - processing done in FPGA



Hash Join : data are in card DDR - processing done in CPU



Action usage

Usage:

```
./snap_hashjoin -usage
Usage: ./snap_hashjoin [-h] [-v, --verbose] [-V, --version]
-C, --card <cardno> can be (0...3)
-t, --timeout <timeout> Timefor for job completion. (default 10 sec)
-Q, --t1-entries <items> Entries in table1. (maximum TABLE1_SIZE defined in $ACTION_ROOT/hw/hw_action_hashjoin.H)
-T, --t2-entries <items> Entries in table2. (maximum TABLE2_SIZE defined in $ACTION_ROOT/hw/hw_action_hashjoin.H)
-s, --seed <seed> Random seed to enable recreation.
-N, --no irq Disable IRQs
```

Example :

```
export SNAP_TRACE=0x0
```

```
cd $SNAP_ROOT && export ACTION_ROOT=$SNAP_ROOT/actions/hls_hashjoin
source snap_path.sh
snap_maint -vv
```

```
echo Random generation of 2 tables with default table size: Table1/Q = 25 entries / Table2/T = 23 entries
snap_hashjoin -vv -t 2500 -C0
echo Random generation of 2 tables with 30 entries for Table1/Q and 60 for Table2/T
=> this will induce 2 calls of the action since Table2 is limited to 32 on purpose
snap_hashjoin -vv -t 2500 -Q 30 -T 60 -C0
```

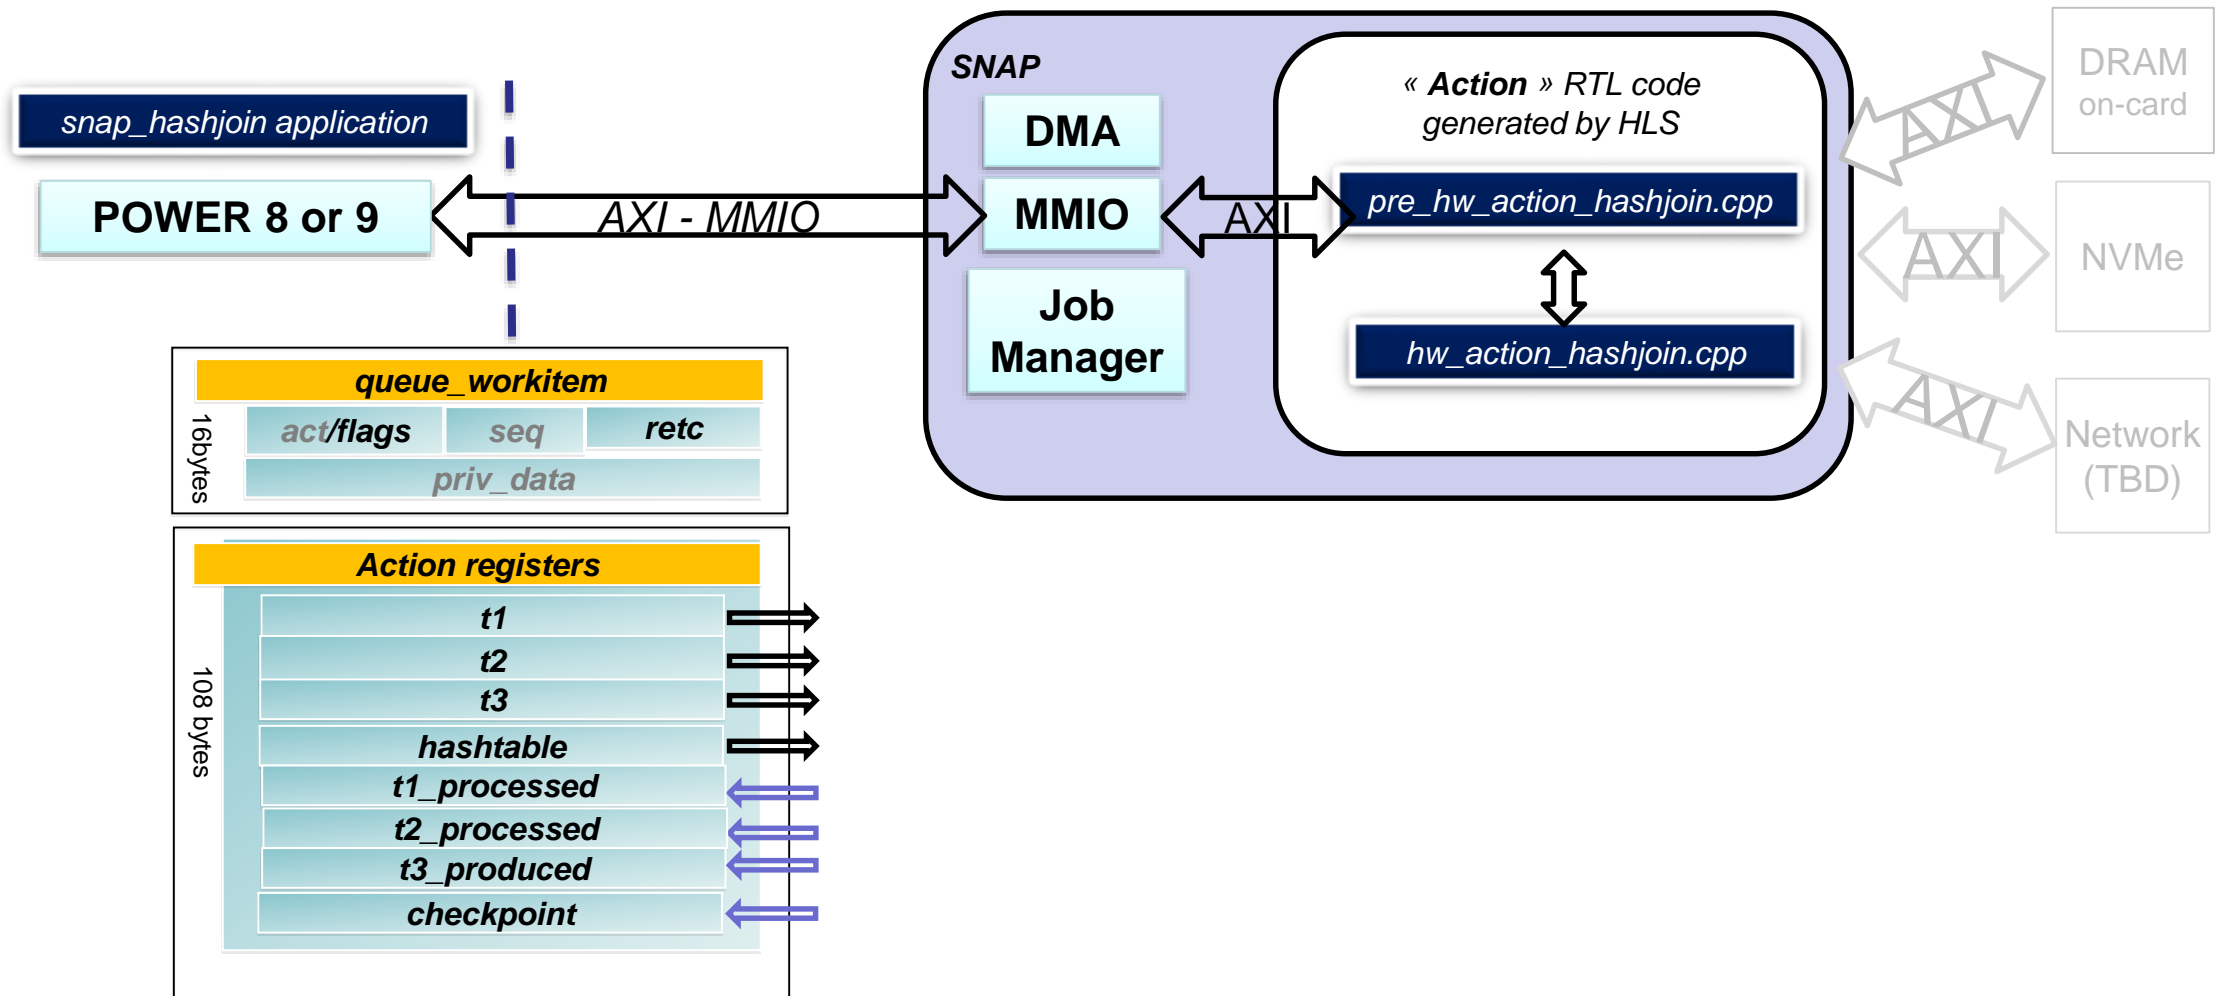
```
echo This example can also be run using no FPGA/ in CPU mode
```

```
SNAP_CONFIG=CPU ./snap_hashjoin -vv -t 2500 -Q 30 -T 60
```

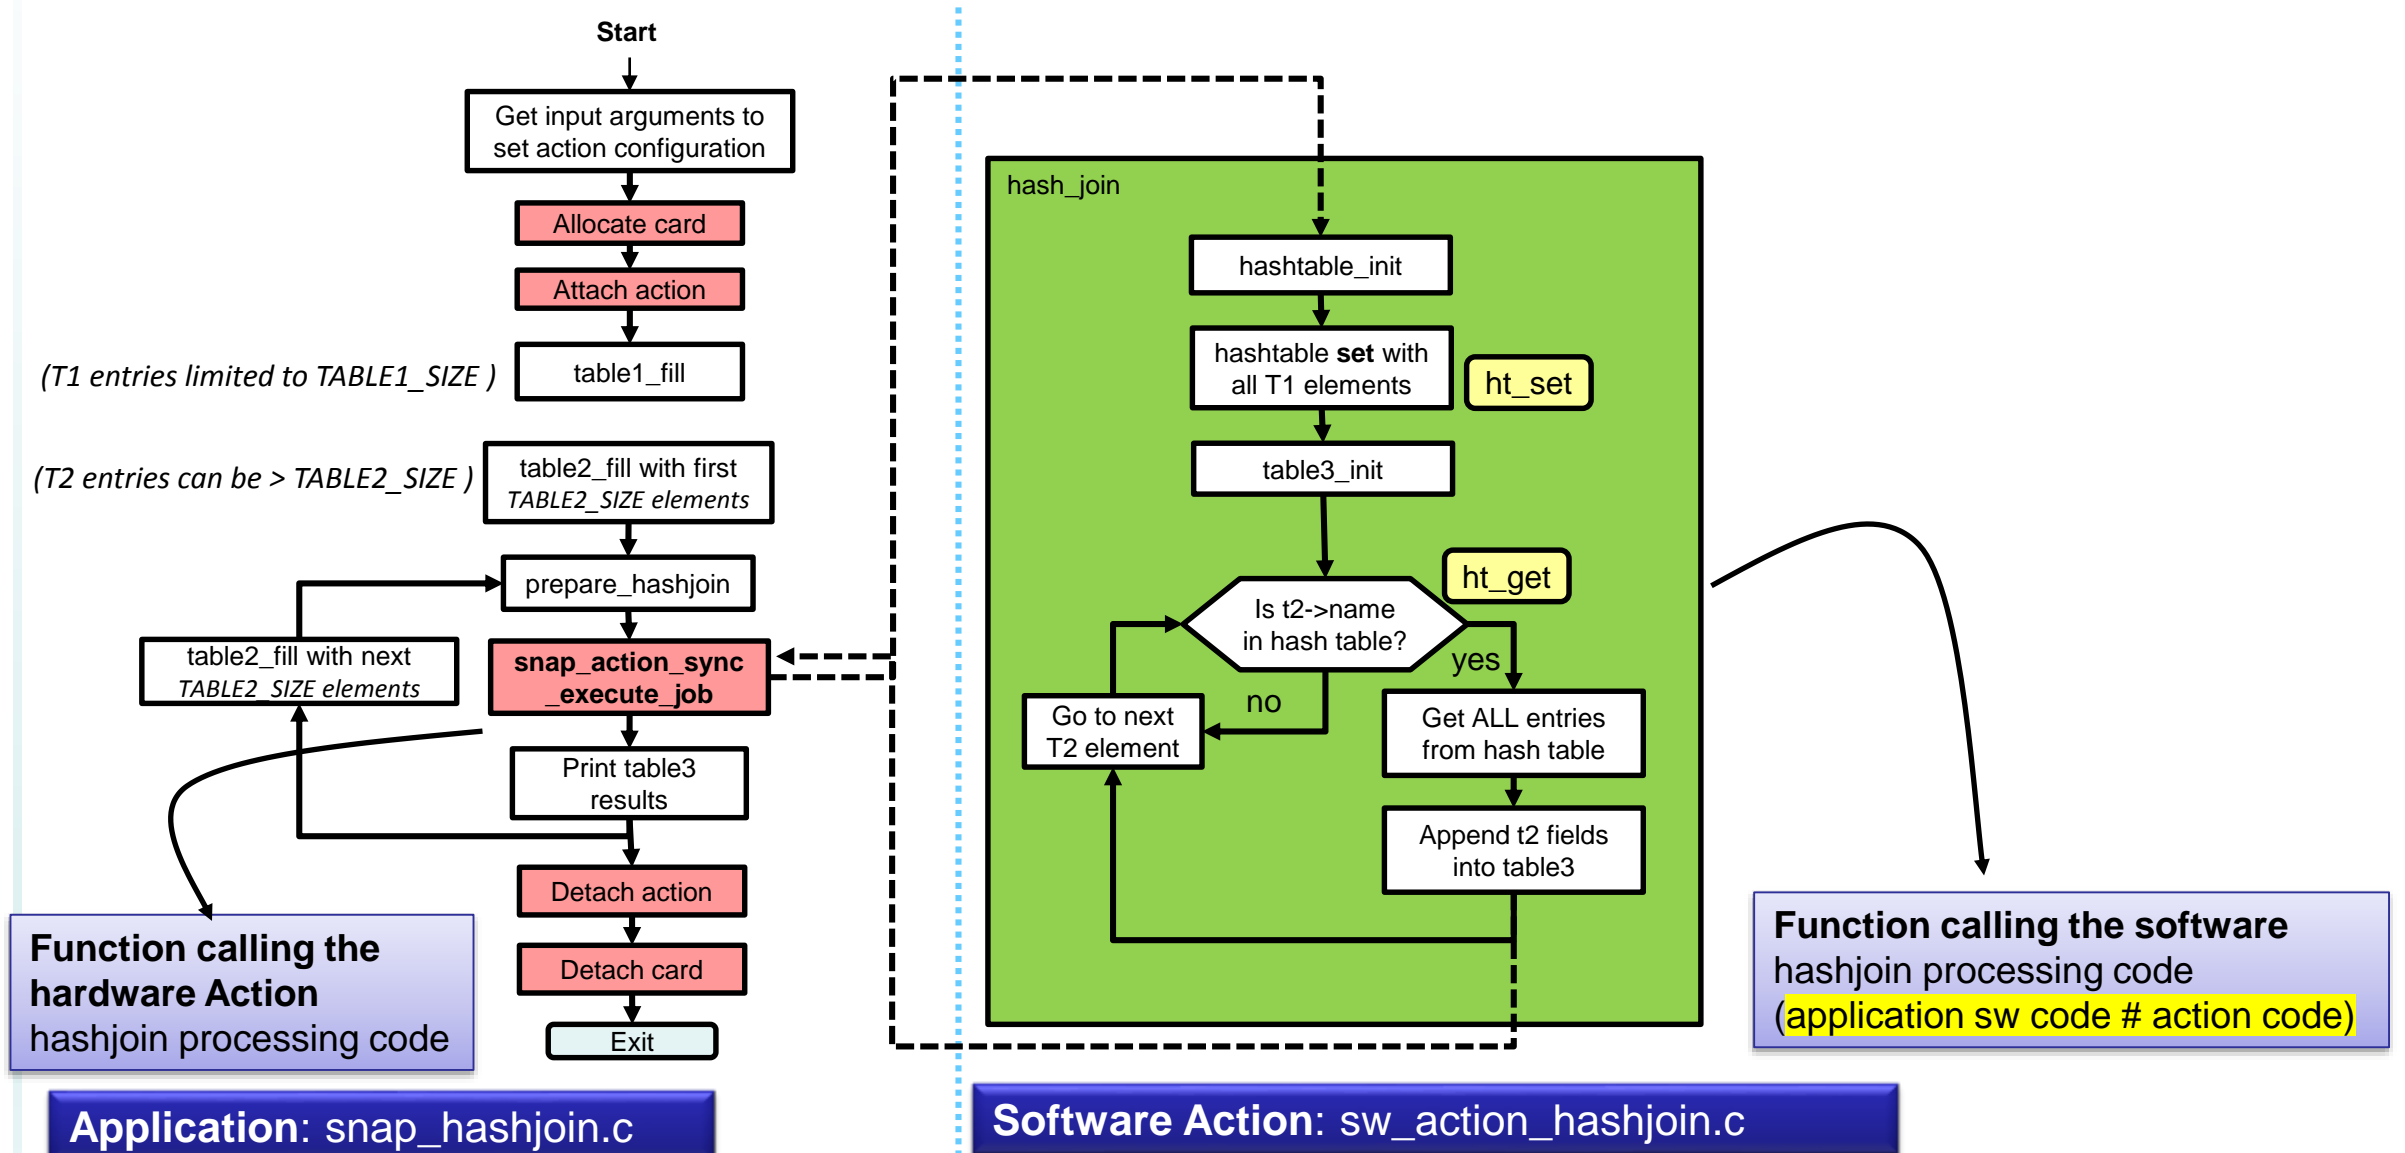
Options:

```
SNAP_TRACE = 0x0 → no debug trace
SNAP_TRACE = 0xF → full debug trace
SNAP_CONFIG = FPGA → hardware execution
SNAP_CONFIG = CPU → software execution
```

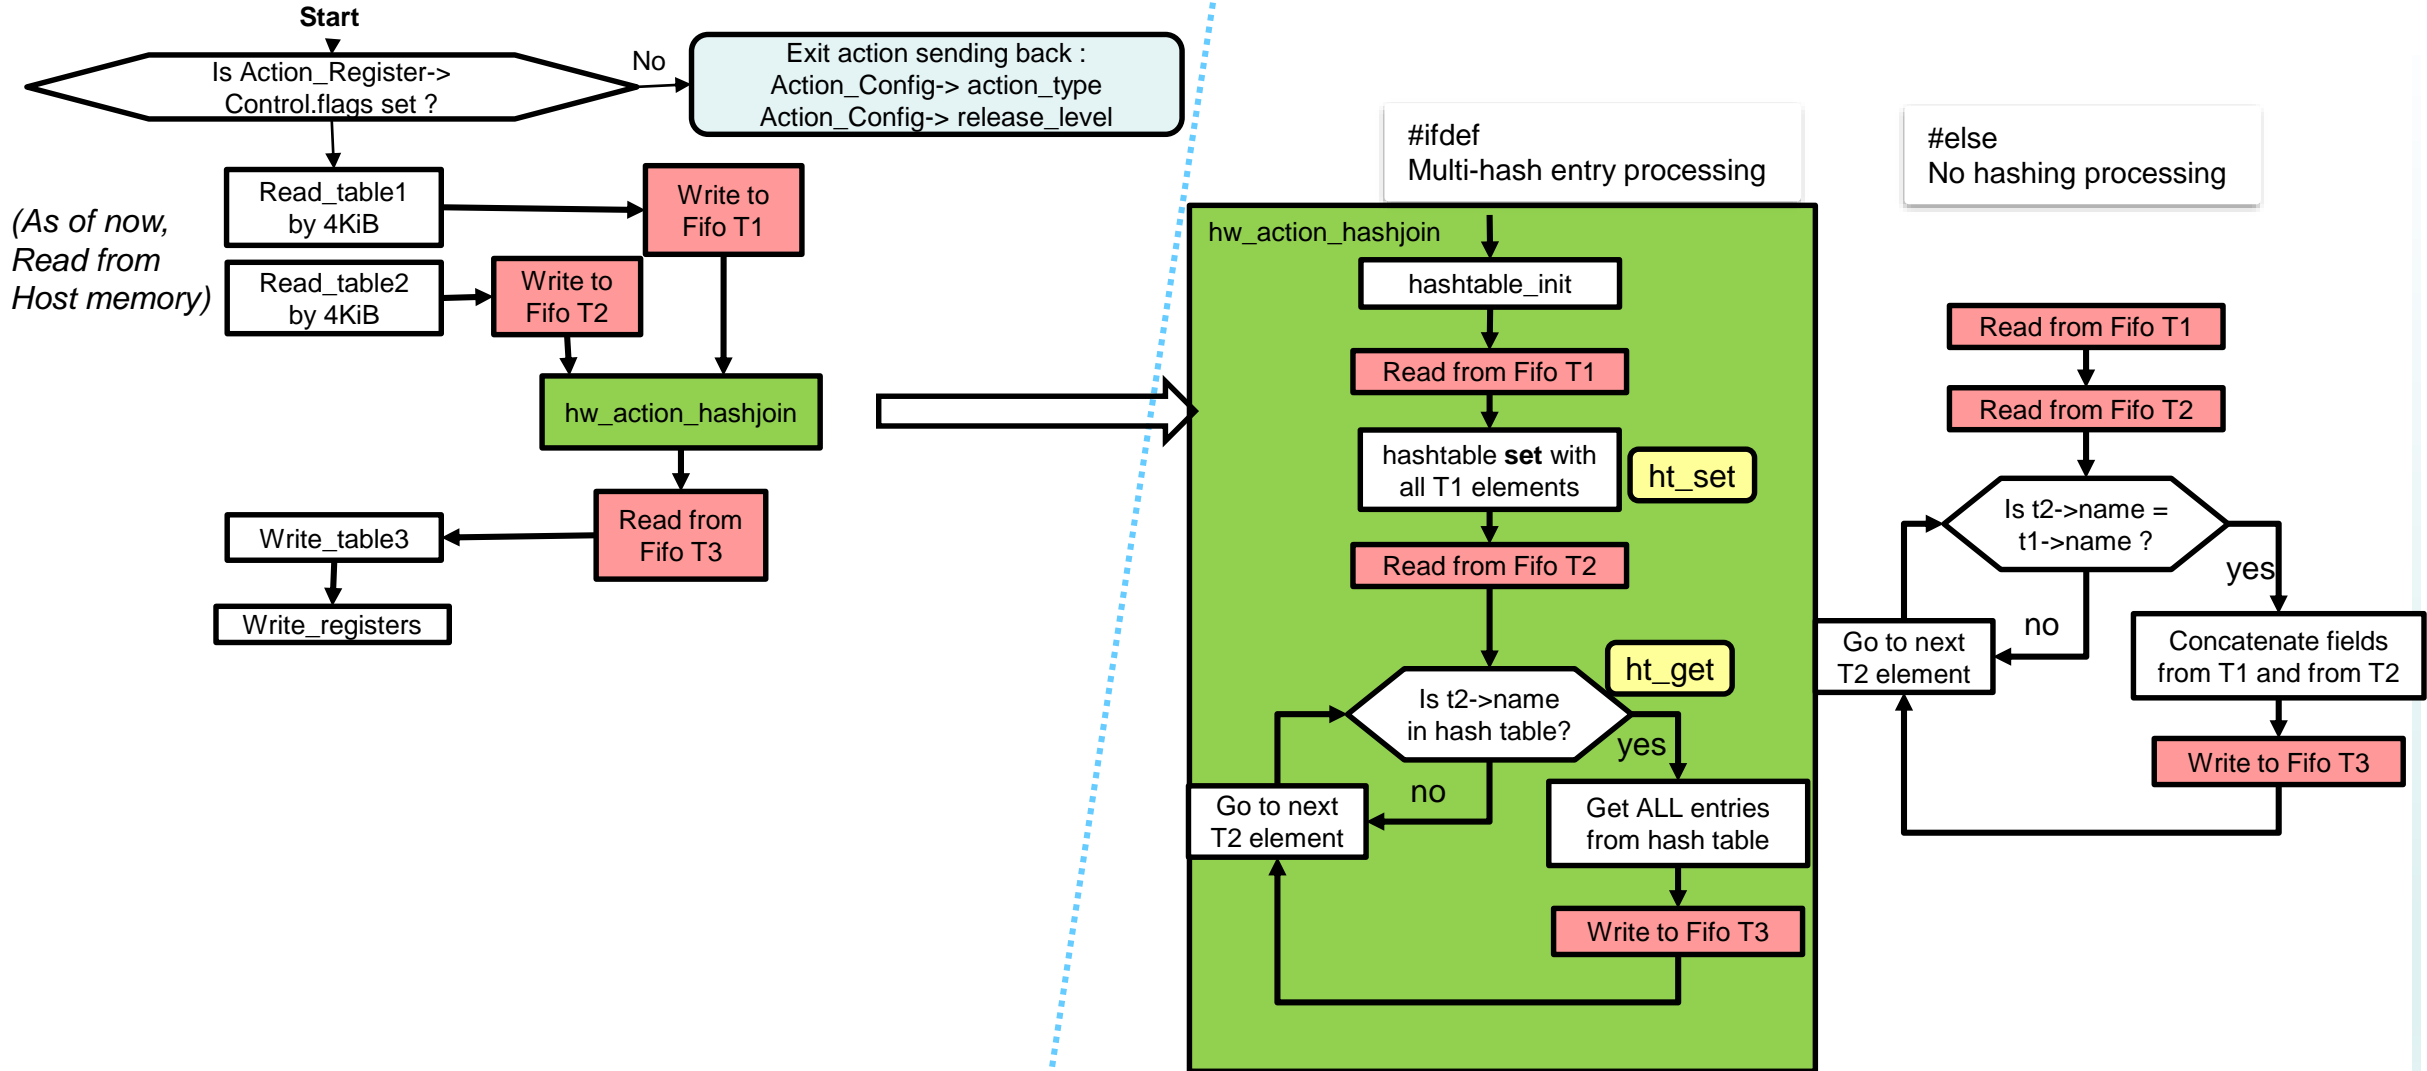

Hashjoin registers



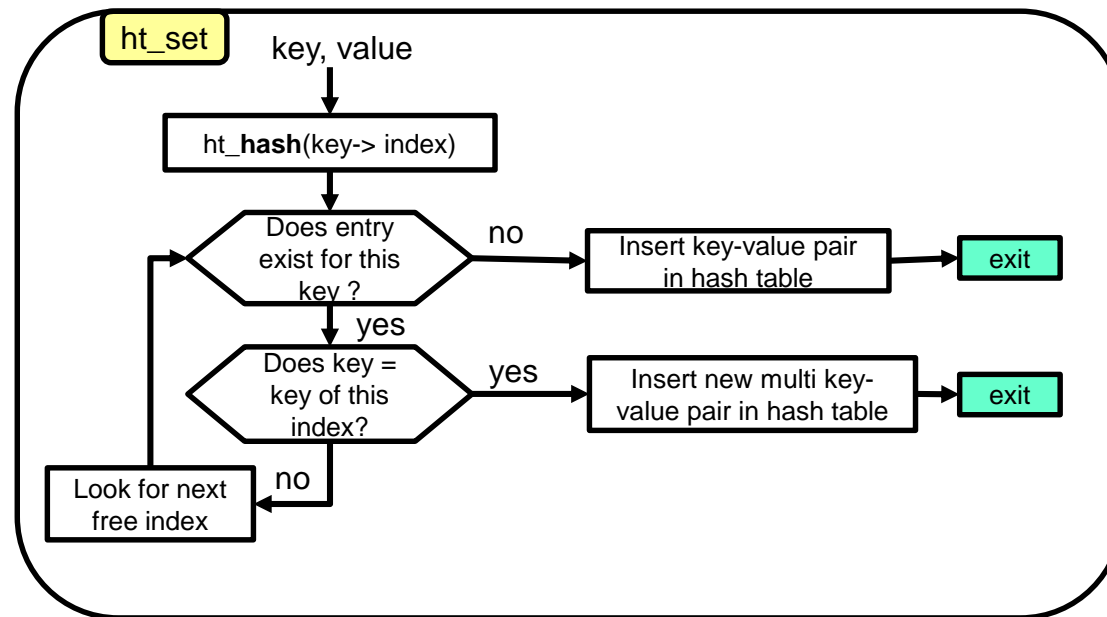
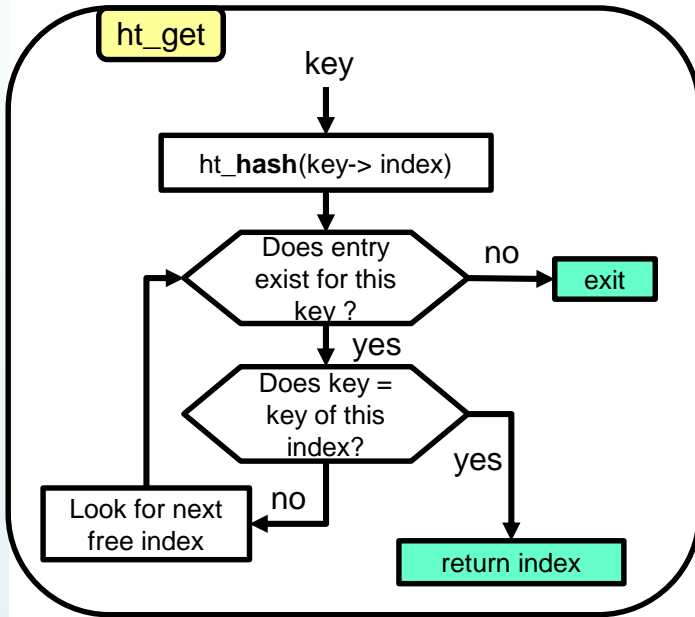
Application Code : what's in it?



Action hashjoin Code : what's in it?



Application - Action hashjoin Code : what's in it?



Common to both → CPU Action : `sw_action_hashjoin.c` → FPGA Action: `hw_action_hashjoin.cpp`

Constants - Ports

Constants: ➔ \$ACTION_ROOT = snap/actions/hls_hashjoin

Constant name	Value	Type	Definition location	Usage
HASHJOIN_ACTION_TYPE	0x10141002	Fixed	\$ACTION_ROOT/include/action_hashjoin.h	Checksum ID - list is in snap/ActionTypes.md
RELEASE_LEVEL	0x00000022	Variable	\$ACTION_ROOT/hw/hw_action_hashjoin.H	release level – user defined
TABLE1_SIZE	32	Variable	\$ACTION_ROOT/hw/hw_action_hashjoin.H	Maximum number of entries for Table1
TABLE2_SIZE	32	Variable	\$ACTION_ROOT/hw/hw_action_hashjoin.H	Maximum size of the Table 2 for the hardware action, but entries can be from any size. Multiple calls to Action will return table3 results
TABLE3_SIZE	(TABLE1_SIZE * TABLE2_SIZE)	Operation	\$ACTION_ROOT/hw/hw_action_hashjoin.H	Table 3 will be from any size
HT_SIZE	(TABLE1_SIZE * 16)	Operation	\$ACTION_ROOT/hw/hw_action_hashjoin.H	Definition of the size of hashtable
HT_MULTI	(TABLE1_SIZE)	Operation	\$ACTION_ROOT/hw/hw_action_hashjoin.H	multihash entries depends on table1

Ports used:

Ports name	Description	Enabled
din_gmem	Host memory data bus input Addr : 64bits - Data : 512bits	Yes
dout_gmem	Host memory data bus output Addr : 64bits - Data : 512bits	Yes
d_ddrmem	DDR3 - DDR4 data bus in/out Addr : 33bits - Data : 512bits	No
nvme	NVMe data bus in/out Addr : 32bits - Data : 32bits	No

MMIO Registers

Read and Write are considered from the application / software side

act_reg.Control		This header is initialized by the SNAP job manager. The action will update the Return code and read the flags value.							
CONTROL		If the flags value is 0, then action sends only the action_RO_config_reg value and exit the action, otherwise it will process the action							
Simu - WR	Write@	Read@	3	2	1	0	Typical Write value	Typical Read value	
0x3C40	0x100	0x180	sequence		flags	short action type	f001_01_00		
0x3C41	0x104	0x184	Retc (return code 0x102/0x104)				0	0x102 - 0x104	SUCCESS/FAILURE
0x3C42	0x108	0x188	Private Data				c0febabe		
0x3C43	0x10C	0x18C	Private Data				deadbeef		
action_reg.Data		Action specific - user defined - need to stay in 108 Bytes							
hashjoin_job_t		This is the way for application and action to exchange information through this set of registers							
Simu - WR	Write@	Read@	3	2	1	0	Typical Write		
0x3C44	0x110	0x190	[snap_addr]t1.addr (LSB)						
0x3C45	0x114	0x194	[snap_addr]t1.addr (MSB)						
0x3C46	0x118	0x198	[snap_addr]t1.size						
0x3C47	0x11C	0x19C	[snap_addr]t1.flags (SRC, DST, ...)		[snap_addr]t1.type (DRAM, NVME,...)				
0x3C48	0x120	0x1A0	[snap_addr]t2.addr (LSB)						
0x3C49	0x124	0x1A4	[snap_addr]t2.addr (MSB)						
0x3C4A	0x128	0x1A8	[snap_addr]t2.size						
0x3C4B	0x12C	0x1AC	[snap_addr]t2.flags (SRC, DST, ...)		[snap_addr]t2.type (DRAM, NVME,...)				
0x3C4C	0x130	0x1B0	[snap_addr]t3.addr (LSB)						
0x3C4D	0x134	0x1B4	[snap_addr]t3.addr (MSB)						
0x3C4E	0x138	0x1B8	[snap_addr]t3.size						
0x3C4F	0x13C	0x1BC	[snap_addr]t3.flags (SRC, DST, ...)		[snap_addr]t3.type (DRAM, NVME,...)				
0x3C50	0x140	0x1C0	[snap_addr]hashtable.addr (LSB)						
0x3C51	0x144	0x1C4	[snap_addr]hashtable.addr (MSB)						
0x3C52	0x148	0x1C8	[snap_addr]hashtable.size						
0x3C53	0x14C	0x1CC	[snap_addr]hashtable.flags (SRC, DST, ...)		[snap_addr]hashtable.type (DRAM, NVME,...)				
0x3C54	0x150	0x1D0	t1_processed						
0x3C55	0x154	0x1D4	t2_processed						
0x3C56	0x158	0x1D8	t3_produced						
0x3C57	0x15C	0x1DC	checkpoint						

\$ACTION_ROOT/hw/hw_action_h

```
typedef struct {  
    CONTROL Control; /* 1  
    hashjoin_job_t Data; /* 1  
    uint8_t padding[SNAP_HLS  
} action_reg;
```

\$ACTION_ROOT/include/act

```
typedef struct hashjoin_job  
    struct snap_addr t1; /*  
    struct snap_addr t2; /*  
    struct snap_addr t3; /*  
    struct snap_addr hash  
  
    uint64_t t1_processed;  
    uint64_t t2_processed;  
    uint64_t t3_produced;  
    uint64_t checkpoint;  
} hashjoin_job_t;
```

\$SNAP_ROOT/actions/include/hls_snap.H

```
typedef struct {  
    snapu8_t sat; // short action type  
    snapu8_t flags;
```

\$SNAP_ROOT

```
typedef str  
    uint64  
    uint32
```

\$ACTION_ROOT/hw/hw_action_hashjoin.H

typedef struct {

CONTROL Control; /* 16 bytes */

hashjoin_job_t Data; /* 108 bytes */

uint8_t padding[SNAP_HLS_JOBSIZE - sizeof(hashjoin_job_t)];

} action_reg;

\$ACTION_ROOT/include/action_hashjoin.h

typedef struct hashjoin_job {

struct snap_addr t1; /* IN: input table1 for multihash */

struct snap_addr t2; /* IN: 2nd table2 to do join with */

struct snap_addr t3; /* OUT: resulting table3 */

struct snap_addr hashtable; /* CACHE: multihash table */

uint64_t t1_processed; /* #entries cached, repeat if not all */

uint64_t t2_processed; /* #entries processed, repeat if not all */

uint64_t t3_produced; /* #entries produced store them away */

uint64_t checkpoint;

} hashjoin_job_t;

\$SNAP_ROOT/actions/include/hls_snap.H

typedef struct {

snapu8_t sat; // short action type

snapu8_t flags;

snapu16_t seq;

snapu32_t Retc;

snapu64_t Reserved; // Priv_data

} CONTROL;

\$SNAP_ROOT/software/include/snap_types.h

typedef struct snap_addr {

uint64_t addr;

uint32_t size;

snap_addrtype_t type; /* DRAM, NVME, ... */

snap_addrflag_t flags; /* SRC, DST, EXT, ... */

} snap_addr_t;

Measured performance

Times are in μs			"SW" hashjoin process (on CPU)			"HW" hashjoin process (on FPGA)						
2x64B		2x64B					1/2 of T2	3x64B			1/2 of T2	1/4 of T2
T1 (entries)	T2 (entries)	T1+T2 Size(Bytes)	DDR to Host	SW	Total	HW	T3 (entries)	Size in Bytes	DDR to Host	Total	Avg Speed up	Avg Speed up
30	50	10,240	7,314	511	7,825	241	25	4,800	3,429	3,670	2.1	4.1
30	500	67,840	48,457	1,014	49,471	1,318	250	48,000	34,286	35,604	1.4	2.7
30	5,000	643,840	459,886	5,132	465,018	12,222	2,500	480,000	342,857	355,079	1.3	2.5
30	50,000	6,403,840	4,574,171	34,493	4,608,664	114,813	25,000	4,800,000	3,428,571	3,543,384	1.3	2.5

For this performance measurement, we considered :

- 2 tables in entries containing each 2 x 64Bytes fields
- 1 table for results containing 3 x 64Bytes fields
- ***We considered that results are 1/2 of the number of the entries of the largest table (realistic ?)***

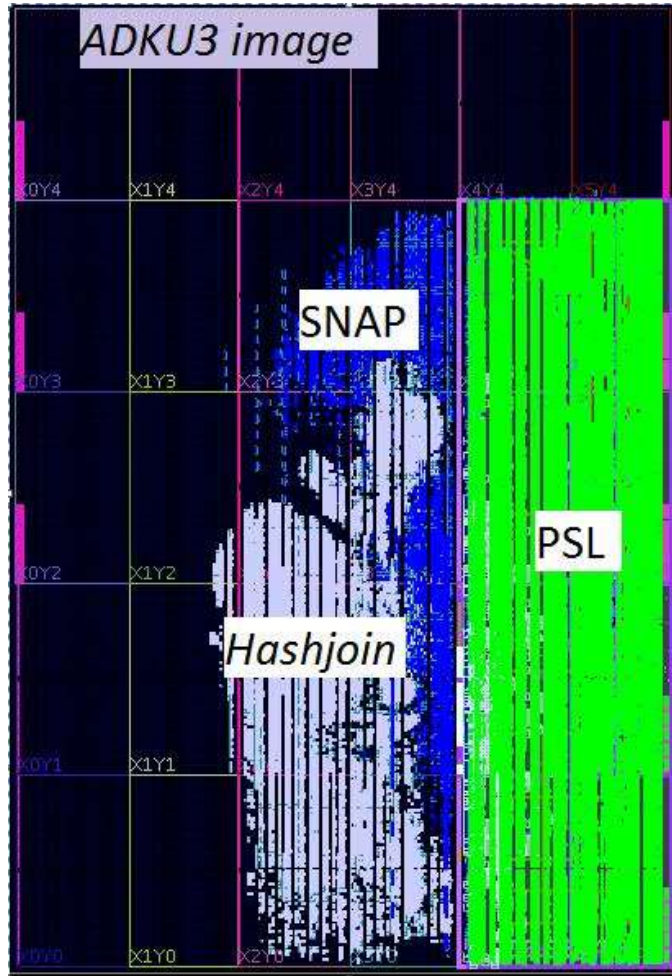
Comment : No real optimization work was done on this example since hash logic may need to be modified to build the image an easier way

What else ?

Path of improvement ?

1. Find how to write the algorithm so that HLS can parallelize :
 - Building hashtable with new elements from table 1 (small table)
 - Checking / adding new elements from table 2 (large table)
 - Filling result table 3➔ This mean handling collision
2. Enable conversion/type-casting of flat memory to structures
3. Support unaligned data access e.g. special FIFOs
4. Pointers, dynamic memory allocation, ...

ADKU3 image of hls_hashjoin (as is, with no optimization)



History of this document and of the action release level

V2.0: initial document

V2.1: new files directory structure applied

V2.2: cleaning code